

Adaptive Tuning of Parallel Programs with CnC

Concurrent Collections (CnC) 2016

Murali Emani



LLNL-PRES-XXXXXX

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under contract DE-AC52-07NA27344. Lawrence Livermore National Security, LLC

 **Lawrence Livermore
National Laboratory**

Computing Landscape is Dynamic

- Parallel systems are continuously evolving
- Diverse architectures, workloads and data


Computing Landscape is Dynamic

- Parallel systems are continuously evolving
- Diverse architectures, workloads and data
- Execution environment is **dynamic** !
- Need adaptive parallelism partitioning and mapping.



Adaptive Parallelism Mapping


- Co-executing programs
- Varying degree of resource contention



Adaptive Parallelism Mapping

Program performance is sensitive to the environment


- Co-executing programs
- Varying degree of resource contention



Adaptive Parallelism Mapping

Tune program for better performance

- Co-executing programs
- Varying degree of resource contention



Tuners

- Is it enough to tune once and leave as-is ?
- Execution environment might change
- Tuning has to adapt to the environment through the program execution

Tuners in CnC

- Several knobs available to tune
- Step priorities
 - Rank steps based on amount / importance of work
`CNC_USE_PRIORITY`
- Thread affinity
 - Improve locality and minimize thread movement
`FIFO_AFFINITY / step_tuner::affinity()`
`CNC_PIN_THREADS`

Tuners in CnC

- Number of threads

`CnC::debug::set_num_threads`
`CNC_NUM_THREADS`

- Tag-Ranges / Tuning Ranges

- Group of steps instead of single instances

- Partitioning ranges

- Hierarchical CnC ?

Current work

(1) Optimal Tuner Selection


Determining


- which parameters to tune and
- how much to tune

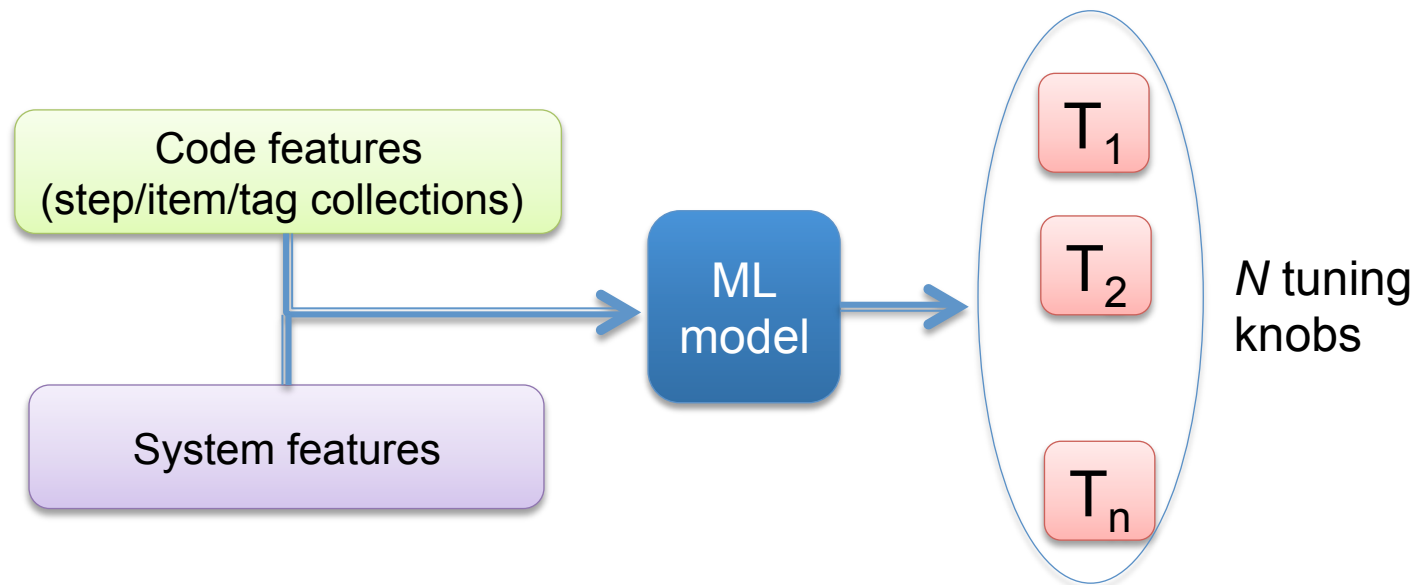
-
- Each tunable parameter may have different impact
 - Strong / Weak Correlation analysis
 - For example,
 - (a) #threads (b) affinity and (c) partitioner
 - Fibonacci series, with different combinations gave up to **2x** speedup !!

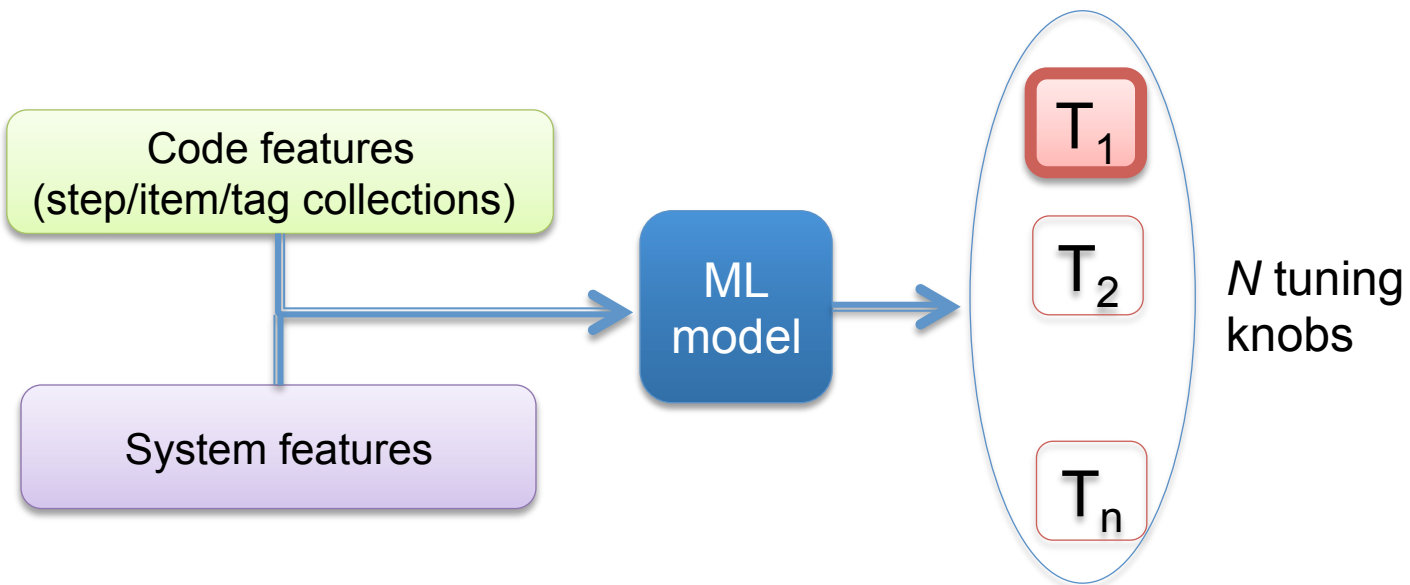
Machine learning could be of help

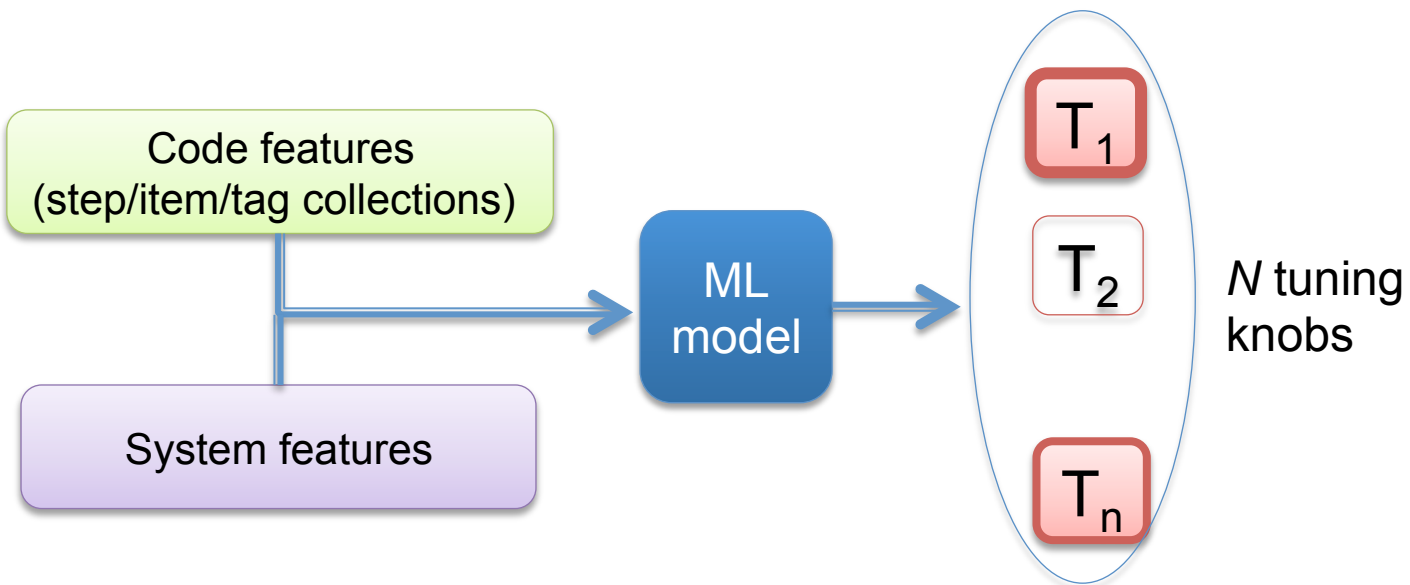
- **ICS '06:** “Online power-performance adaptation of multi-threaded programs using hardware event-based prediction”, Curtis-Maury, M., Dzierwa, J., Antonopoulos, C. D., and Nikolopoulos, D. S.
- **ISPASS'12:** “Using Utility Prediction Models to Dynamically Choose Program Thread Counts.”, Moore, R. W. and Childers, B. R
- **PLDI'12:** “Parcae: A System for Flexible Parallel Execution”, A. Raman, A. Zaks, J. W. Lee, and D. I. August.
- **CGO '13:** “Smart, Adaptive Mapping of Parallelism in the Presence of External Workload”, M. Emani, Z. Wang and M. O'Boyle
- **PLDI'14:** “Adaptive, Efficient, Parallel Execution of Parallel Programs”, Sridharan, G. Gupta, and G. S. Sohi.
- **PLDI'15:** “Celebrating Diversity: A Mixture of Experts Approach for Runtime Mapping in Dynamic Environments”, M. Emani and M. O'Boyle
- **LCPC '16:** “Mapping Medley: Adaptive Parallelism Mapping with Varying Optimization Goals”, M.Emani













(2) How to determine optimal #threads



Step collection

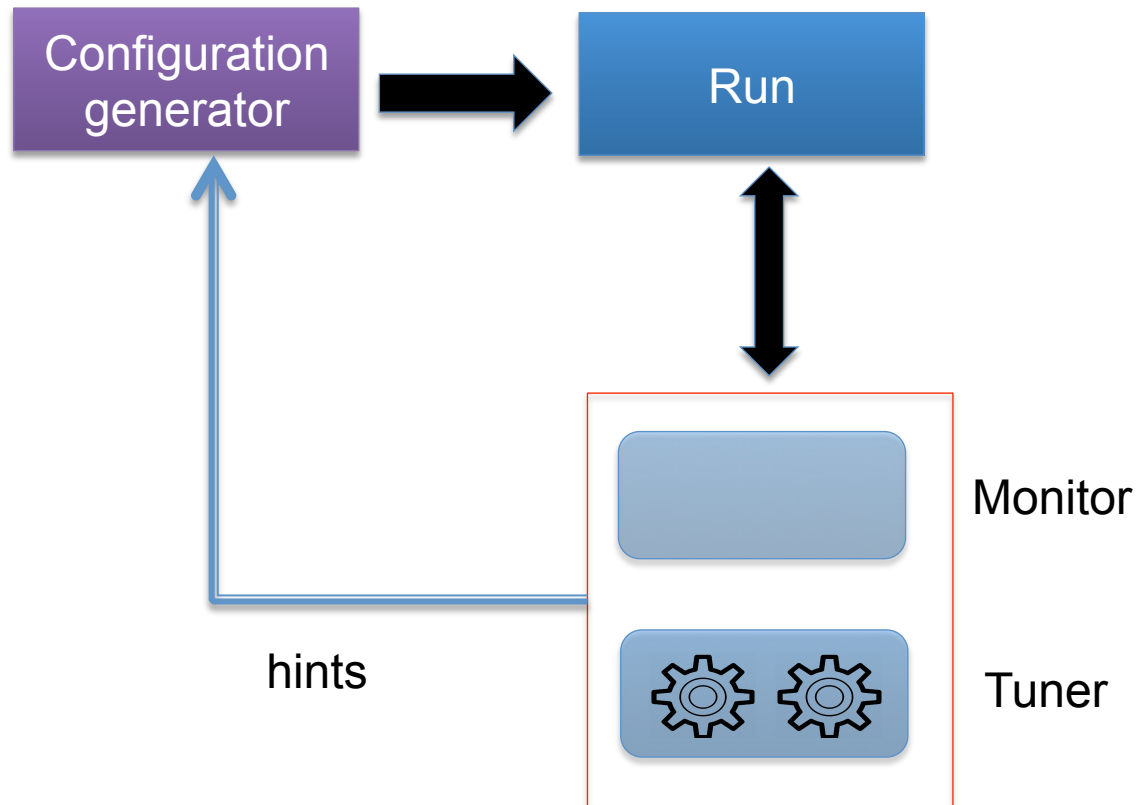


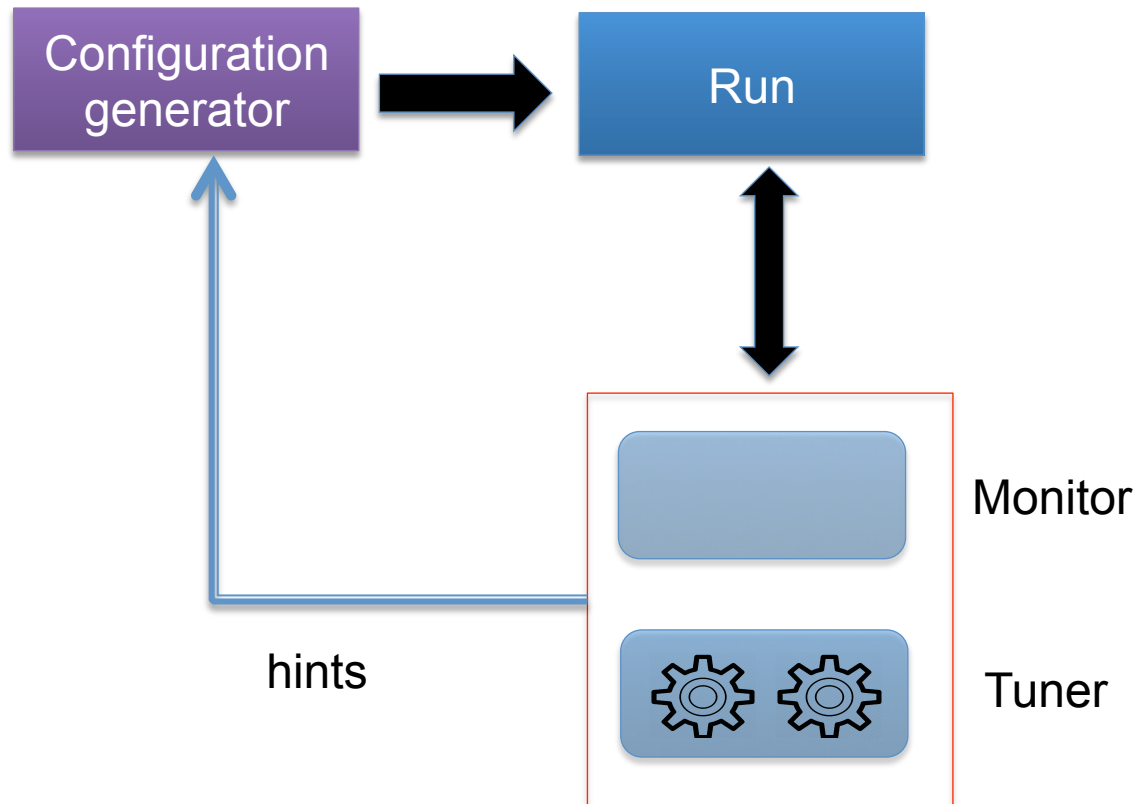
Schedule all instances



Schedule range
of instances

How it works

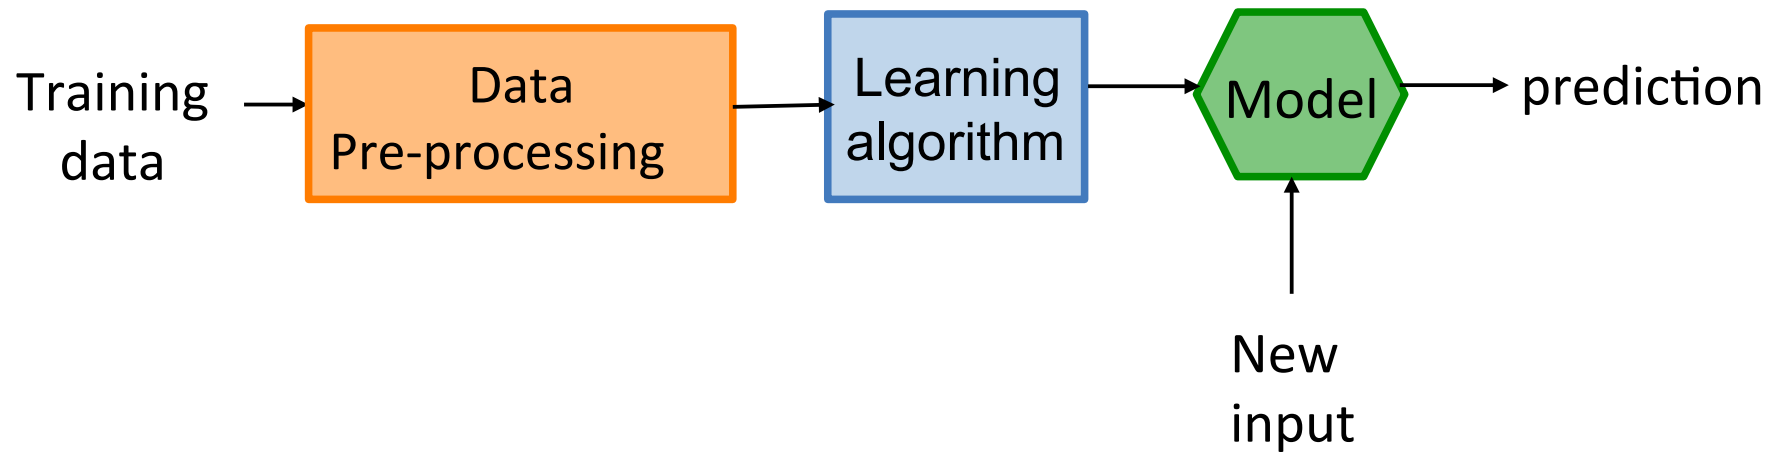




- May take time to reach optimal configuration
- Machine Learning can help reach quicker

Approach – Machine Learning

- Hand crafted solutions infeasible



- Train offline, deploy online

Challenges

- Training Data
 - Experiment Design space
 - Prune data
- Supervised / Unsupervised
 - Semi-supervised ?
 - Start with a learned heuristic,
 - update and re-learn as and when required

-
- Active / Passive invoke
 - *Heartbeat* mode – monitor at regular intervals, or
 - As and when required
 - Ensure no additional overhead

Summary

- Tuning needs to be adaptive
- Faster decision to tune can deeply impact performance
- Machine Learning could be one way forward.

Murali Emani,
Lawrence Livermore National Laboratory,
emani1@llnl.gov