

EVALUATING PERFORMANCE OF TASK AND DATA COARSENING IN CONCURRENT COLLECTIONS

Chenyang Liu, Milind Kulkarni

Overview

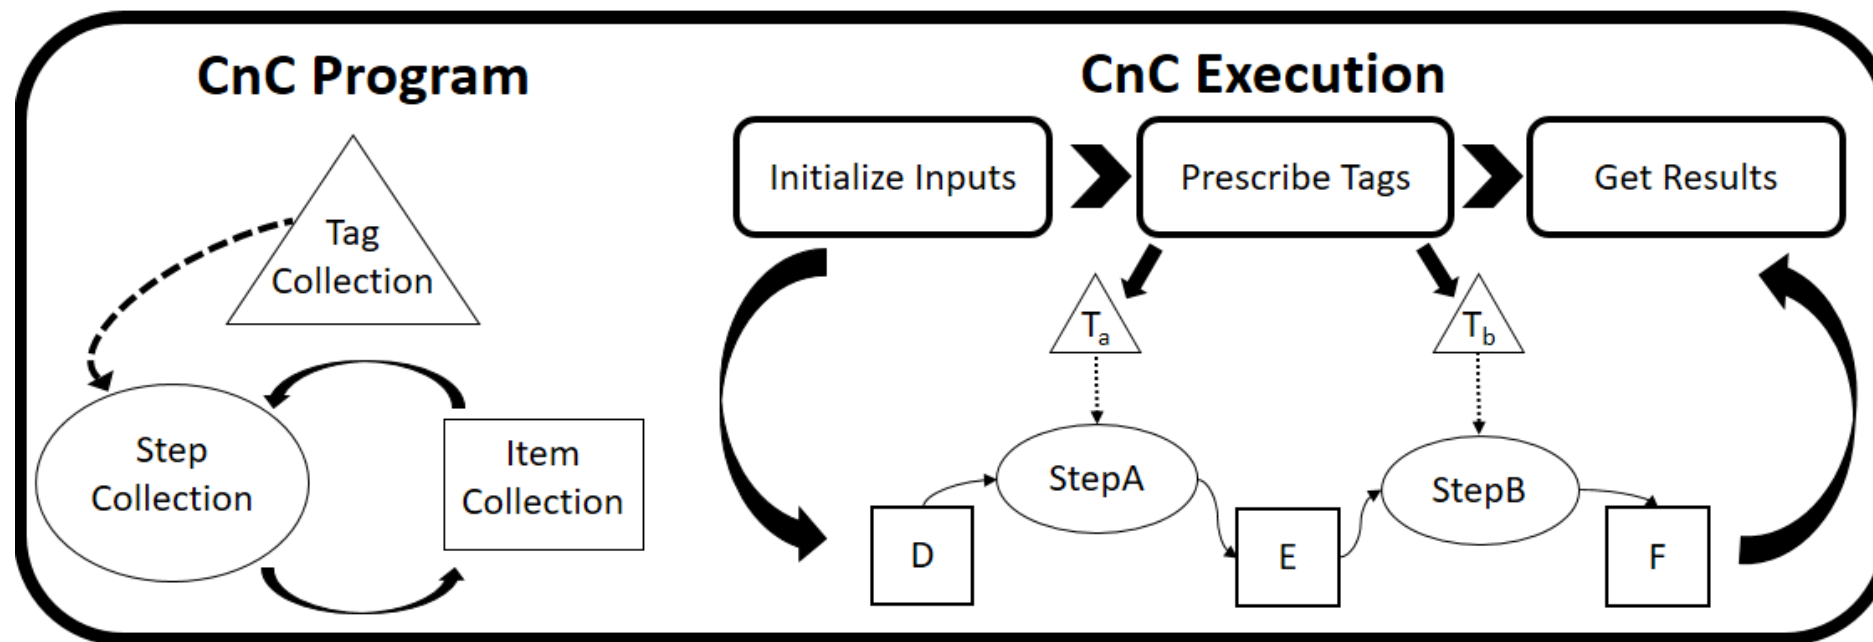
- Scientific Applications
 - Domain expert required for correctness, algorithm semantics, etc.
 - Performance expert for tuning performance, scalability, portability, etc.
- Concurrent Collections
 - “Separation of Concerns” Philosophy - decouple performance and algorithmic/domain concerns
- Contributions
 - Evaluate Performance benefits
 - Discuss High-Level transformations
 - Explore Automation and Tuning

Scientific Applications

- Programmability
 - Algorithmic correctness - requires understanding of the scientific domain or method
 - Algorithmic design – modular data/functions (ease of future programming)
 - Exploit semantic optimization (high-level tuning)
- Performance
 - Parallelization challenges
 - Expressing Parallelism
 - CPUs? GPUs? NUMA?
 - Memory Hierarchy/Locality
 - Communication Costs
- *Idea: Explore high-level optimizations that will improve both programmability and performance for complex parallel applications*

Concurrent Collections

- A Parallel Programming Philosophy
 - “Separation of the Concerns” for domain expert and performance tuning expert
 - Flexible task-parallel model/runtime supports multiple platforms
 - Parallelism automatically exploited using task-based data-driven model
- High-level specifications, Low-level tuners
 - Declarative specification for computation/data/control dependencies
 - Hardware tuners allow machine-specific optimization
- Research
 - Task-level fusion/tiling
 - Data-tiling
 - Programming automation




- **Data-Driven Model**
 - Tags prescribe steps, creating dynamic step instances
 - Steps execute when inputs are ready (step-like property)
- **Data is immutable (Dynamic single assignment)**
 - Key/Value lookup using “get/put” operations
- **Parallel Runtime**
 - Exploits parallelism given dependency constraints

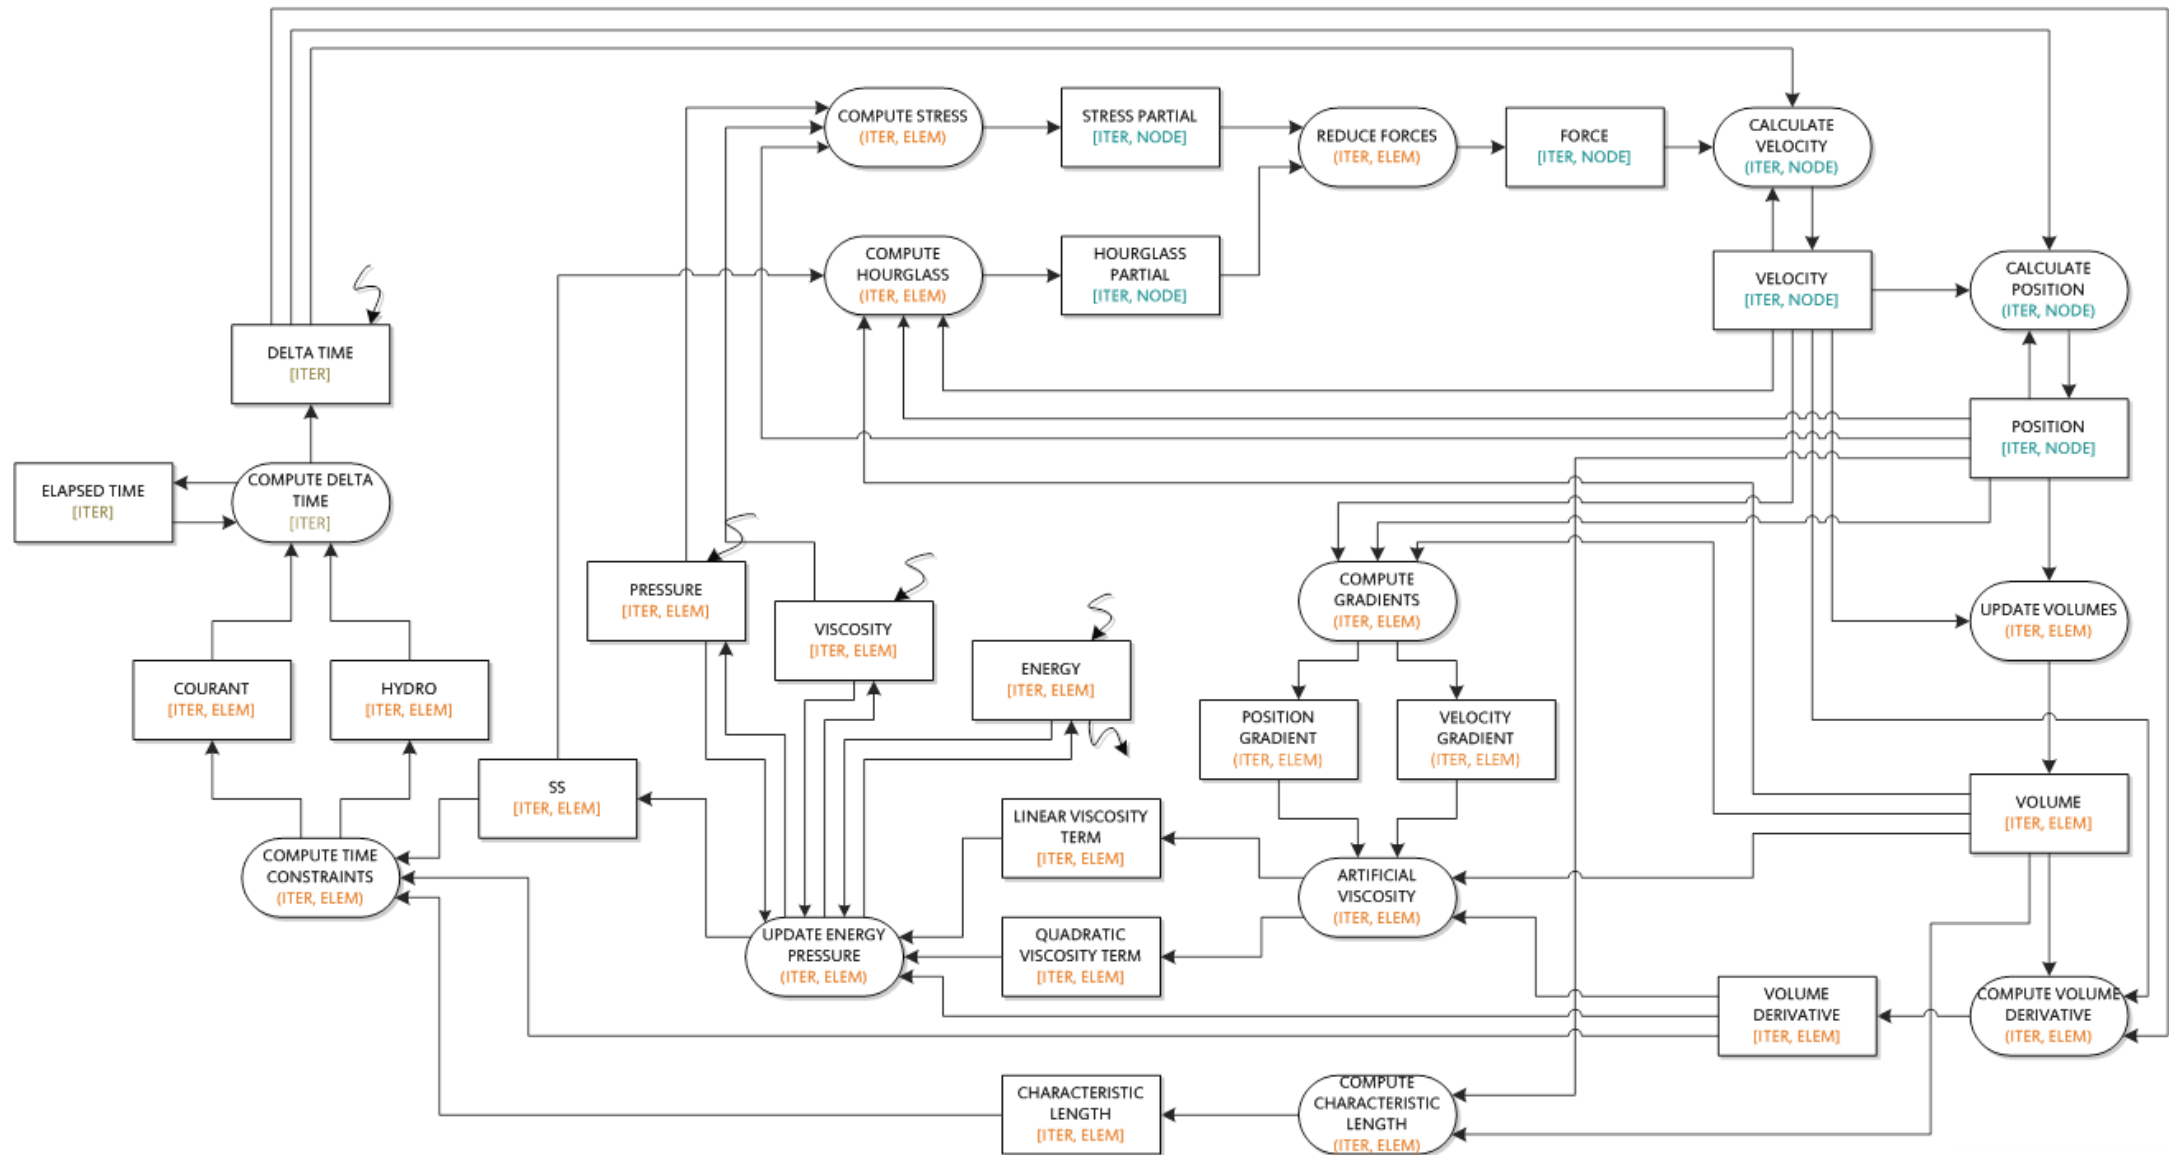
Problem

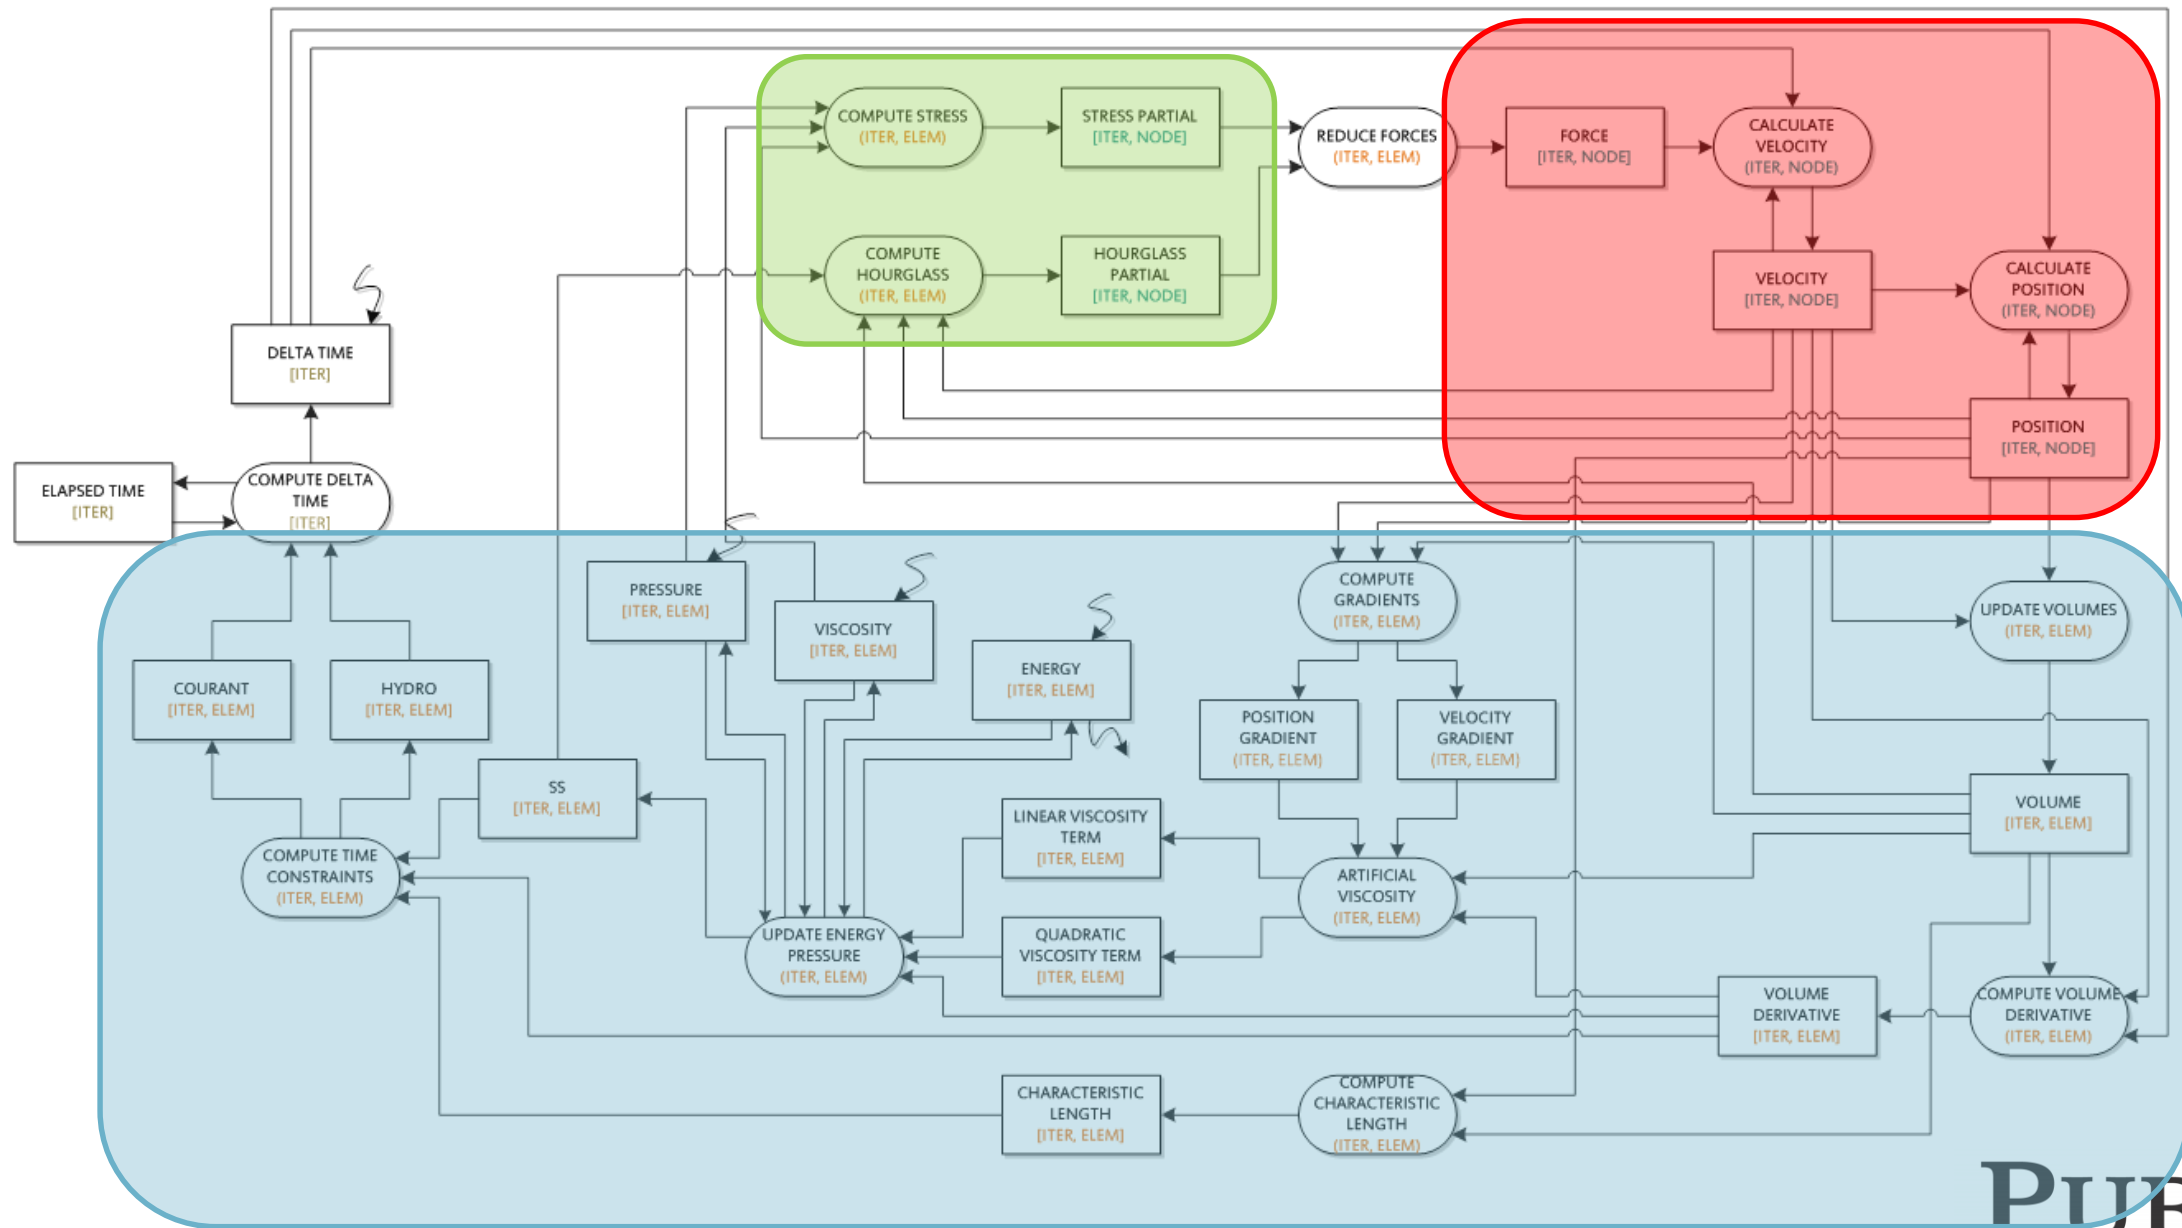
- How to optimize parallel performance for the LULESH code?
 - Start with decomposed algorithm, explore high-level transformation
 - Later modify data layout and create tiled computation steps
- Challenges
 - Task Granularity: Fine-grain parallelism
 - Data Granularity: Synchronization for large datasets
 - Explore Task and Data coarsening (Tiling)
- Other Concerns
 - CnC Translator for semi-automatic code generation
 - Application or machine-specific tuning

CnC LULESH

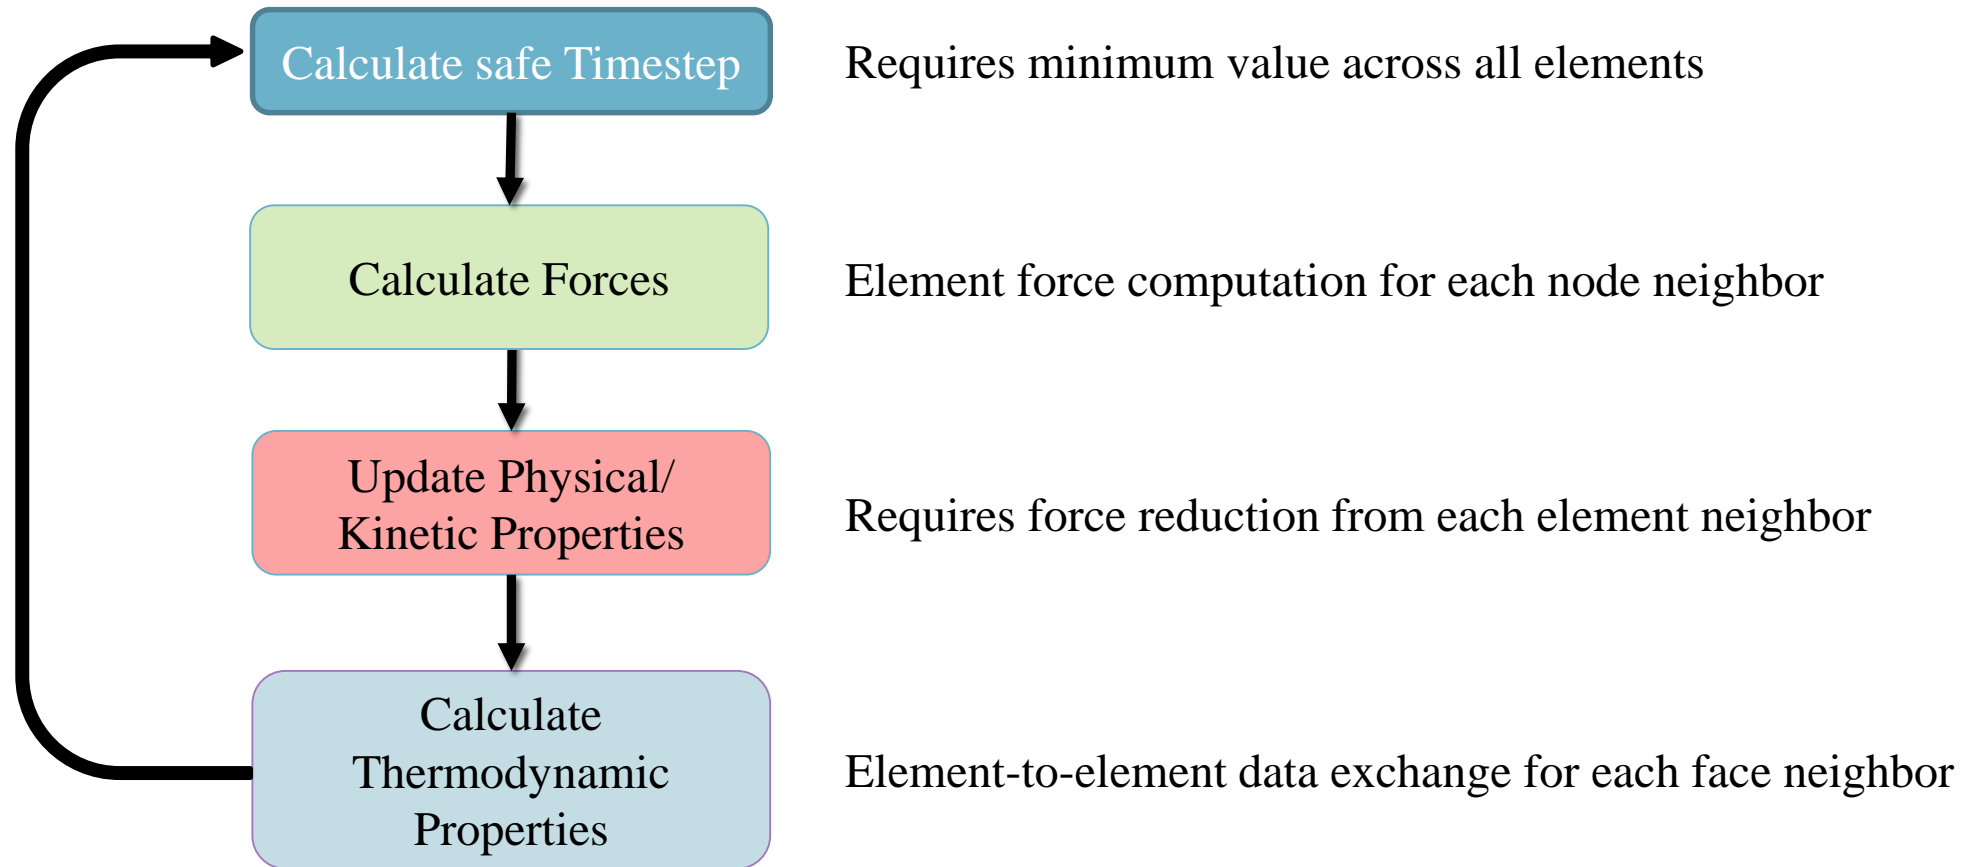
- **LULESH: Livermore Unstructured Lagrangian Explicit Shock Hydrodynamics**
 - Challenge problem from the DARPA UHPC program
- 3D blast wave propagation simulation
 - Operates on a hexahedral mesh with 2 centerings:
 - Node/Element interactions/computations
 - Lots of control-flow
 - Multiple stencil updates
 - Ample Parallelism
- Build upon a fully decomposed algorithm







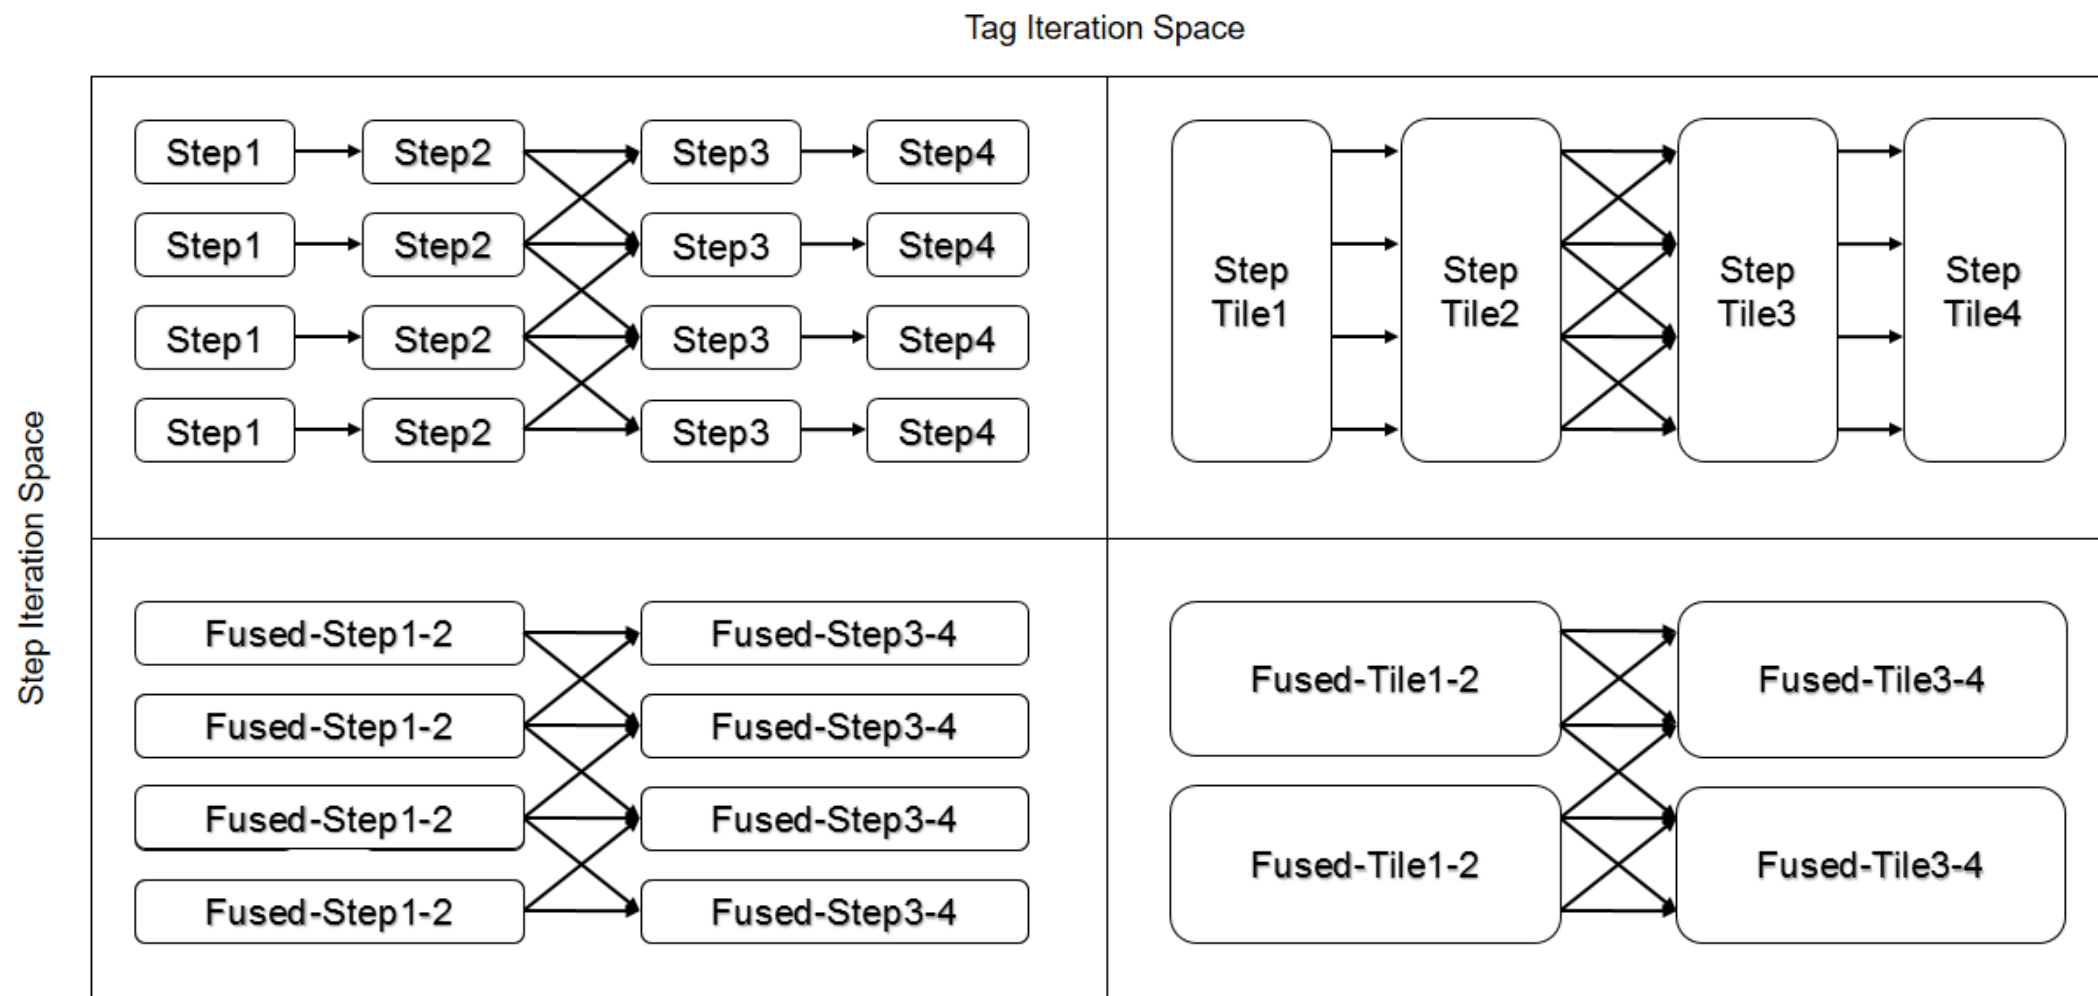
LULESH Algorithm



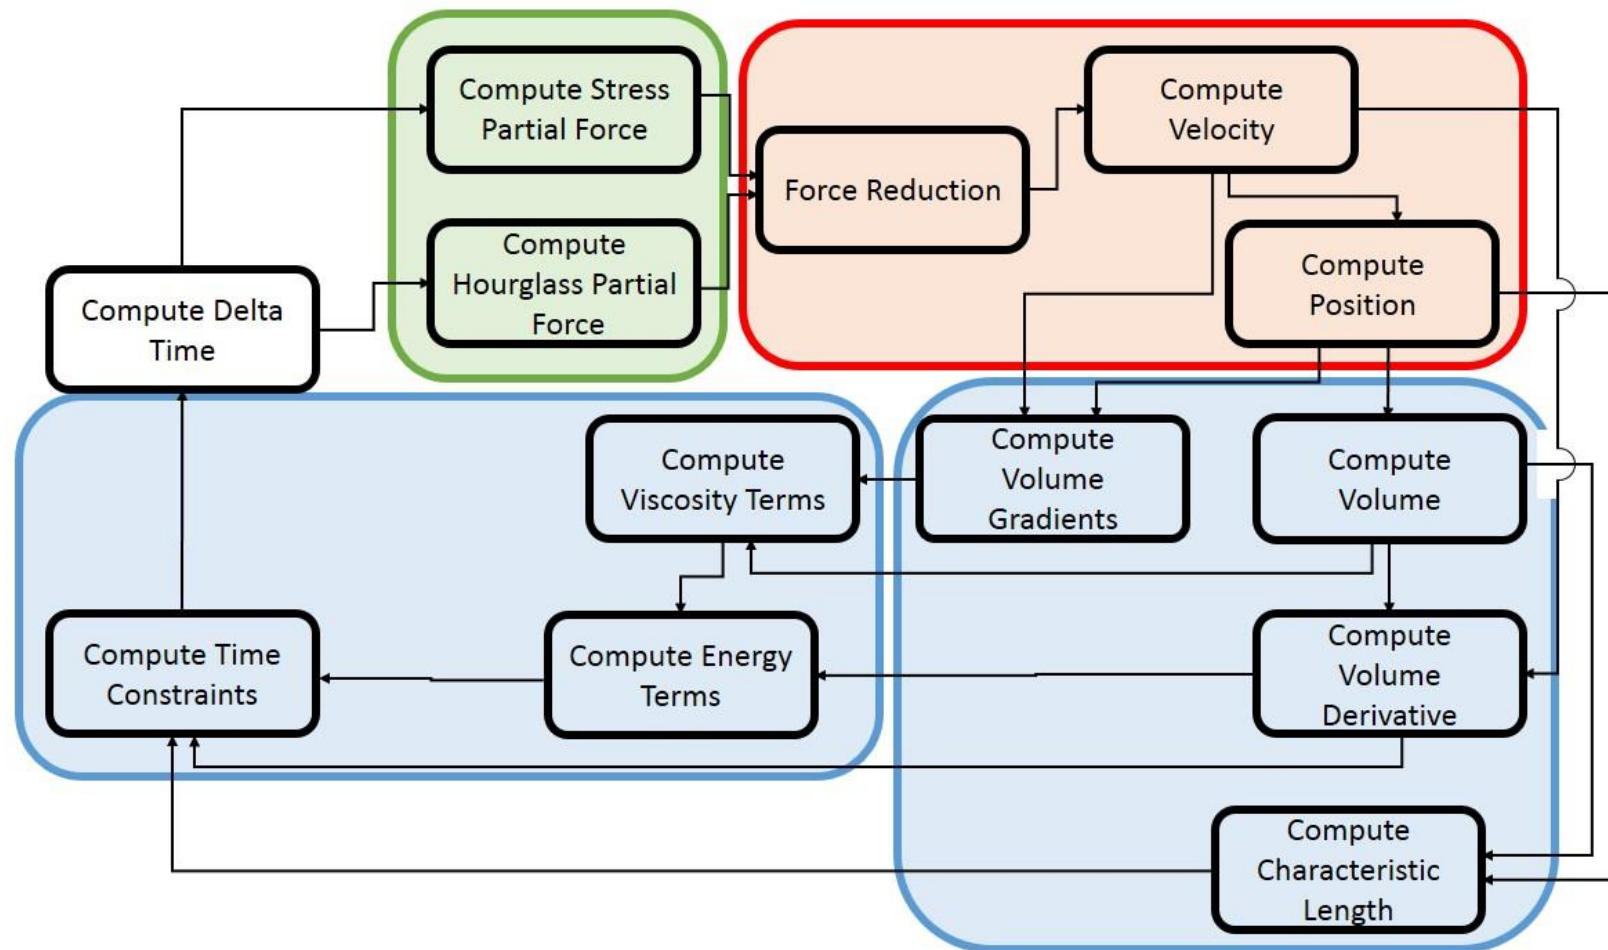
Task Coarsening

- Baseline program specifies the computation per iteration, per node/element
 - Too much overhead from fine-grain parallelism
- Solution: Coarsen through modifications through Collections
- Step Fusion
 - Serialize different steps operating under the same tag
 - Legal as long as no dependency cycles (co-routine)
- Tag Tiling
 - Serialize the same steps operating under different tags
 - Resulting tiled steps must be “step-like”

Step Fusion vs Tag Tiling



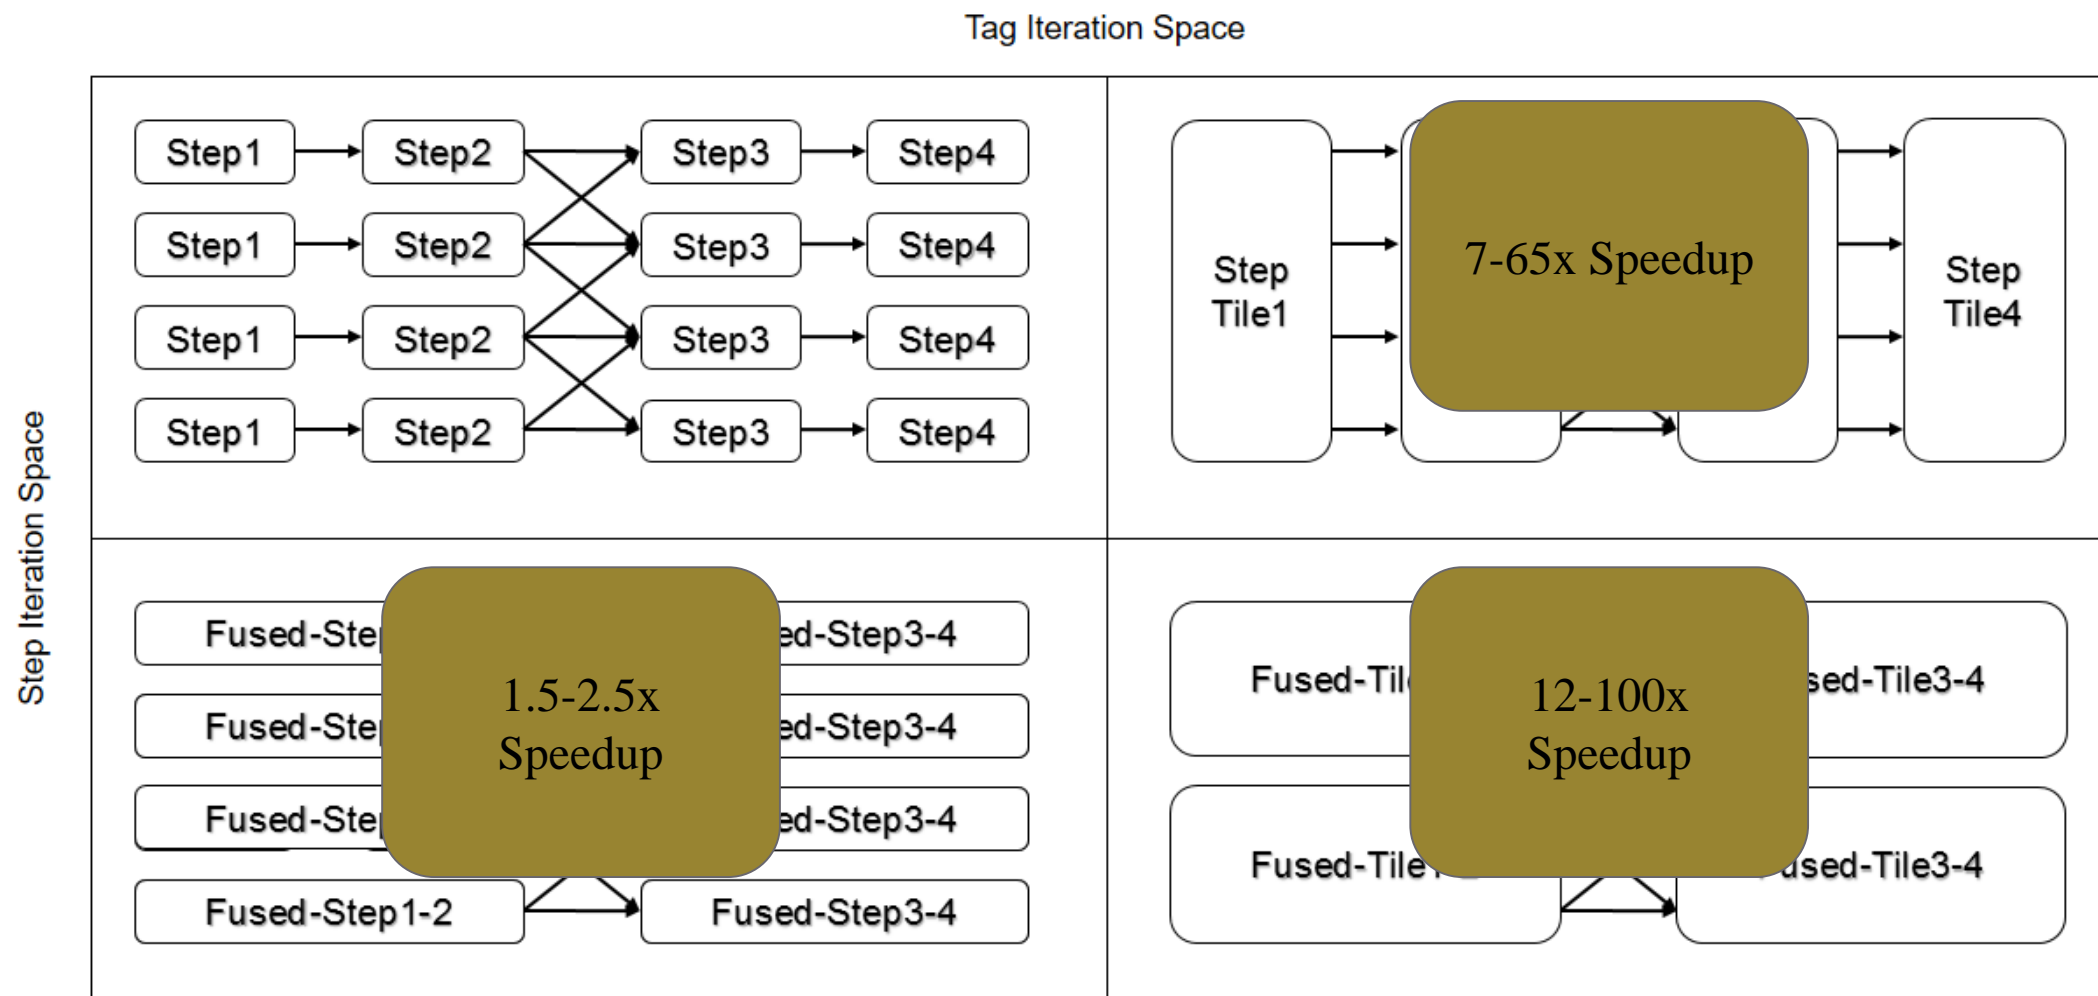
LULESH: Fused Algorithm



Step Fusion & Tag Tiling


- **Challenges:**
- Structural changes to collections and tag/step organization
 - Semi-automatic CnC Translator
- Fused steps require aggregated dependencies
 - Still need to maintain step-like behavior ('get's first, puts after)
- Tiled steps need intermediate storage for temp. data
 - Possible re-use/sharing of neighboring data (code modifications)

Step Fusion vs Tag Tiling




Data Tiling


- Coarsened tasks, but data layout unchanged
 - Locality within tasks, but runtime synchronization required for data in item collection
- Challenges: How to tile inconsistent (node/element) centerings
 - Number of nodes = elements + 1
 - Recall stencil operation that requires update via neighbors
 - Data tiles must conform to dynamic-single-assignment
- Underlying code modifications required for data tiling
 - Kernel computations operates on tiled datasets
 - Further locality/re-use optimizations (variable elimination)
 - Modified communication patterns – intra-tile communication (fast)
- *Note: Data coarsening opportunity a result of task coarsening*




1-Element Force
Computation




2-Nodal Force
Reduction/Update





3-Element Update/
Computation






1-Element Force
Computation







1-Element Force
Computation





2-Nodal Force
Reduction/Update






2-Nodal Force
Reduction/Update







2-Nodal Force
Reduction/Update




⊗ Redundant Interface Node




2-Nodal Force
Reduction/Update




Updated Interface Node




3-Element Update/
Computation





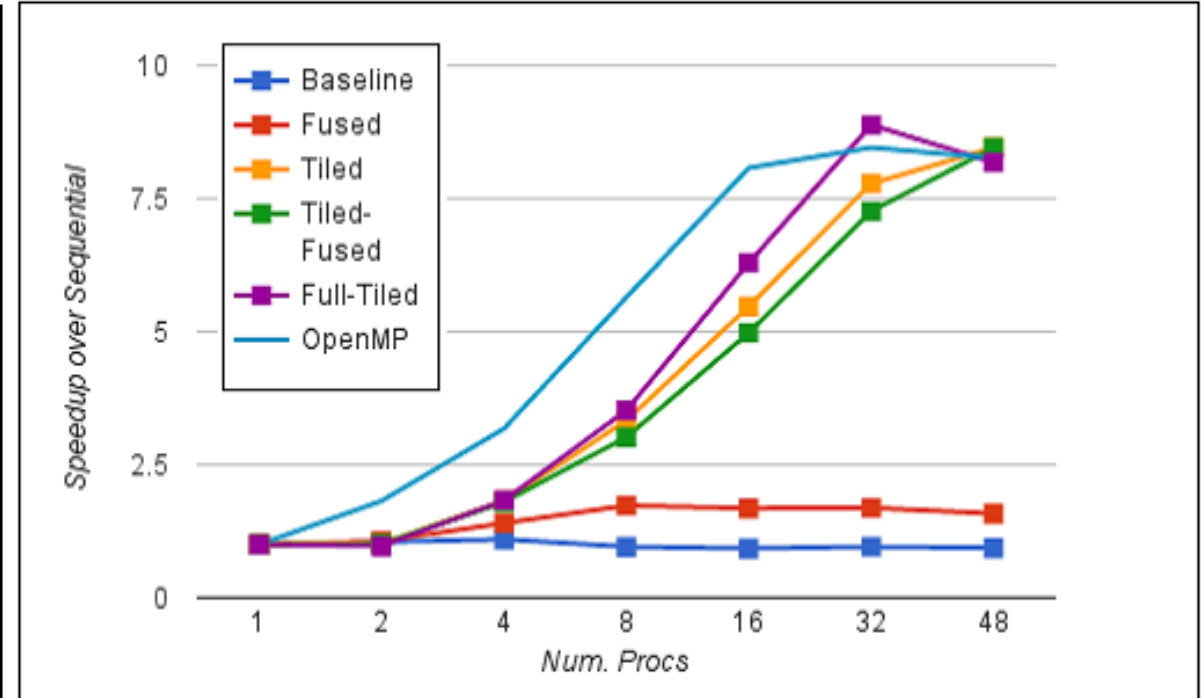
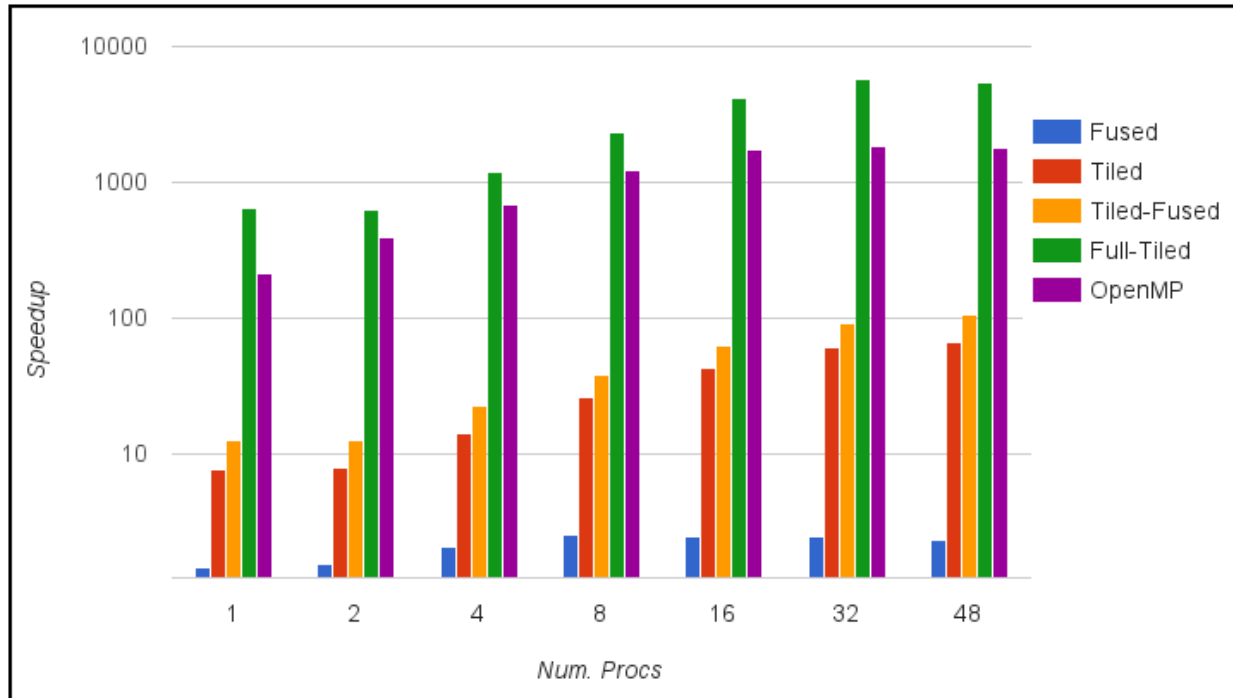
After element computation,
there is an element gradient update:
requires face neighbors



Experimental Results - LULESH

- AMD Opteron 6176 SE system with four 12-core processors (48 cores total) running at 2.3 GHz.
- gcc 4.7, -o3, Intel CnC implementation
- Experiments: Mesh sized 60^3
 - Baseline
 - Fused-only
 - Tiled-only
 - Fused+Tiled
 - (Full) Data Tiled
 - OpenMP (LULESH 2.0.3 LLNL)

Speedup and Scalability



- OpenMP outperforms non-data-tiled implementations by ~10x
- [Full] Data Tiled is 3x faster than OpenMP
- Tag Tiling required for scalability, better than OpenMP (dedicated scheduling process)

Contributions/Discussion

- Task coarsening helps obtain scalability
 - Legal transformations for step fusion & tag tiling at a high-level
- Task+Data coarsening gives strong parallel performance
 - Requires more underlying code modularity/changes
- Ease of Application/Programmability
 - CnC Translator: generates code skeleton from CnC graph (Nick)
 - Programmer: Still required to write sequential code - never worry about parallelism
 - Difficulty: Optimizing data layout and blocked routines for kernel computation
 - Required domain knowledge to capture semantic properties of method

Related

- Related Work

- Various CnC Papers
- Task Parallel Model – Charm++, OpenMP, Chapel, Legion
- Fusion and Tiling – “*Fusion of Parallel Array Operations*”, MRB Kristensen

- Mentions

- Nick Vrvilo – “*Declarative Tuning for Locality in Parallel Programs*”
- Ellen Porter – HPC applications using CnC
- Kath Knobe, Zoran Budimlic - Discussions

Conclusion

- CnC offers an easy way to obtain performance with ease of programmability
 - Provides the right abstractions for domain scientists and performance experts to decouple concerns
 - Gives a good platform to explore high-level optimizations such as computation reordering/tiling
 - Further performance improvements likely requires machine-specific tuning (HW mapping, etc)
- Future Improvements
 - Enhance CnC translator for automatic code generation
 - Pair data templates for 2D/3D spatial containers for common applications
 - Automatically locate redundant data synchronization for performance improvement

Thanks!

Concurrent Collections cont.

- Collections: General Program Structure
- Step Collections
 - Stateless computation tasks
- Item Collections
 - Data input/outputs used by steps
 - Explicit dependencies
 - Dynamically-single assignment using hashmap
- Tag Collections
 - Identifiers for the step collections
 - Dictates program control flow

```

struct lulesh_context:public
    context<lulesh_context>{

    // Step Collections
    step_collection<compute_dt>
        step_compute_dt;
    step_collection<reduce_force>
        step_reduce_force;
    ...

    // Item Collections
    // per node items
    item_collection<pair,vector>force;
    item_collection<pair,vector>position;
    item_collection<pair,vector>velocity;
    // per element items
    ...

    // Tag Collections
    tag_collection<pair>iteration_node;
    tag_collection<pair>iteration_element;
    tag_collection<int>iteration;
    ...

    // Producer Dependencies
    step_compute_dt.consumes(dt);
    ...

    // Consumer Dependencies
    step_compute_dt.produces(dt);
    ...
  }

```

Declaration of CnC Specification